

AI Backend Engineer (1+ YOY) building production RAG systems and agentic workflows in healthcare and SaaS. Shipped a post call intelligence pipeline processing 200+ calls/day and a multi-tenant document intelligence platform for regulatory filings. Work primarily in Python and Node.js with LangGraph, FastAPI, and OpenAI.

TECHNICAL SKILLS

Languages: Python, JavaScript, Java
Backend: FastAPI, Node.js (Express)
Frontend: React, HTML, CSS, TypeScript
Databases: MongoDB, PostgreSQL, Pinecone
AI Systems: RAG pipelines, LangGraph, LangSmith, OpenAI API, Cohere
Infra & Tools: Redis, Docker, Git, REST APIs, JWT

Work Experience

Reinvent Digital

AI Automation Expert

October 2025 – Present

Project: Marketing Attribution & Automation Platform

- Built and integrated a healthcare marketing attribution platform tracking patient interactions across calls, WhatsApp, forms, and ads into GHL CRM, implementing first-touch and last-touch attribution logic across channels
- Built an event-driven automation engine using n8n + WhatsApp API to handle missed-call workflows, reducing median lead response time from 4 hours to under 8 minutes.
- Built a post-call intelligence pipeline (STT → LLM summarization → lead scoring) processing 200+ calls/day; validated on sampled Hindi/English call reviews (95% accuracy), Tamil/Telugu (85% accuracy)

Tech stack: Node.js, Express.js, React, MongoDB, REST APIs, WhatsApp Business API, Google Ads API, Google Business Profile API, n8n, Sarvam Speech-to-Text, LLM API, Docker.

Excollo

Software Engineer

Feb 2025 – October 2025

Project: AI-Persona

- Architected a multi-tenant AI-Persona SaaS backend with full tenant isolation, RBAC, and audit logging serving multiple concurrent workspaces using Node.js, Express, PostgreSQL, and Prisma.
- Engineered automated ingestion pipelines (n8n + OpenAI) processing company documents into Pinecone vectors, reducing manual data-prep effort by 80%.
- Designed a hybrid BM25 + dense vector retrieval pipeline with Cohere reranking to improve answer relevance (validated via internal testing)

Project: CRUDO - Order Management System

- Engineered a multi-tenant Order Management System backend (Node.js, MongoDB) with ERP integration and Pinecone-powered customer clustering supporting multiple concurrent workspaces with consistent query performance
-

PROJECTS

Multi-Agent Document Intelligence System (AI-powered Contract Analysis)

April 2026

- Built a **multi-agent AI pipeline (LangGraph + GPT-4o)** for automated contract review
- Implemented **real-time streaming (SSE)** and **LLM cost optimization via prompt caching (~50% reduction)**
- Used **FastAPI + Next.js** with structured outputs (Pydantic) and **LangSmith tracing for observability**

Technologies used: LangGraph, OpenAI GPT-4o, FastAPI, Next.js, React, TypeScript, Pydantic, LangSmith, SSE

Stock Research Agent (LangGraph)

Feb 2026

- Built a 5-agent stock research system using LangGraph with parallel fan-out (news, financials, sentiment running concurrently) delivering Buy/Hold/Sell reports for US and Indian markets in 30s
- Instrumented full LLM observability via LangSmith tracked per-agent latency and cost-per-report at \$0.0003/run
- Implemented state merging and aggregation logic across parallel agents using LangGraph reducers

Technologies used: Python, LangGraph, LangSmith, FastAPI, OpenAI, Tavily, yfinance

AI-Powered Regulatory Document Intelligence Platform (RHP/DRHP)

Oct 2025 - Jan 2026

- Built a **production RAG system** for regulatory PDFs (300+ pages) with recursive chunking, structured prompts, and **<3s query latency**, reducing manual review effort
- Designed a **multi-tenant backend** with workspace isolation, RBAC, and JWT auth for concurrent clients
- Automated **ingestion** → **embedding** → **summarization** → **comparison pipelines** using n8n, cutting processing time by **~80%**

Technologies used: Node.js, Express, MongoDB, React, TypeScript, n8n, OpenAI APIs, Pinecone, Cohere

EDUCATION

MCA | UPES

2023-2025

BCA | St.Xaviers College

2020-2023